

Table 1: Evaluation results of agents on EXPLAINBENCH, using gpt-5-nano as the question-answering LLM to calculate explanation score.

Agent	ExplainBench				SWE-bench	
	End-to-End		Local		Expl. Score	Efficacy
	<i>Intent</i>	<i>Effect</i>	<i>Intent</i>	<i>Effect</i>		
refact	0.762	0.756	0.157	0.314	0.497 (#2)	0.778 (#3)
Lingxi	0.746	0.762	0.174	0.291	0.493 (#3)	0.785 (#2)
OpenHands	0.757	0.735	0.175	0.344	0.502 (#1)	0.727 (#4)
trae-agent	0.690	0.714	0.168	0.321	0.473 (#4)	0.818 (#1)
mini-SWE-agent	0.523	0.519	0.183	0.247	0.368 (#5)	0.599 (#5)

Supplementary material for “ExplainBench: Evaluating Code Explanations from Agents”

GPT-5-nano ranking results

Results of running EXPLAINBENCH using gpt-5-nano as the answering LLM are presented in Table 1. Note that the agent ranking based on explanations is identical even while the answering LLM has changed, suggesting our ranking is resilient to choice of LLM.